

Permuting the surrogate signal

Research note: is dL/dp_i better than its own random permutation?

2026-09-27.

1 Purpose and summary of conclusions

The RBLapSum backward is not the derivative of its hard forward, which admits the suspicion that the surrogate gradients it produces are “wrong” – that the score path works through its scale and structure rather than through the task signal it routes. This note tests that suspicion with a matched-marginals control. A fraction ρ of each token’s Top- $(K+J)$ pool is sampled every training backward, and the $dL/dp_i = u_i z_i$ values at those positions are shuffled by a uniform random permutation of the subset – *before* the kernel weighting, so every position keeps its own κ_i , the κ -weighted zero-sum correction applies to the permuted signal, and the hard task path is untouched. The construction preserves the surrogate’s scale profile, kernel locality, zero-sum structure and the per-token multiset of signal values; it destroys only the *assignment* of signal to neuron. $\rho = 0$ is bit-for-bit the unmodified backward (`rblapsum_rho_random_perm_prob_grad`, training only).

Conclusions:

- **The assignment is the signal.** Final CE degrades monotonically in ρ at both completed windows: $+0.038/ +0.076/ +0.292$ nats at $(K, J) = (32, 224)$ and $+0.012/ +0.066/ +0.236$ at $(32, 480)$ for $\rho = 0.25/0.5/1$. Even a quarter of the pool permuted costs a measurable amount from the first thousand steps, and the gap never closes.
- **Fully permuted signal is worse than no surrogate at all.** At $\rho = 1$ the runs end at 1.7925 and 1.7269, clearly above the plain hard Top-32 baseline (1.6162, no surrogate, no output norm). Correctly scaled, correctly localized, zero-sum noise delivered through the full surrogate machinery actively damages training; it does not merely waste the term.
- **Narrow windows destabilize under full permutation.** The $\rho = 1$ cells at $J = 32$ and $J = 96$ were at CE 7.7 (step 1500) and 6.0 (step 2500) and climbing when the campaign was stopped, while the same ρ survives 20k steps at $J = 224$ and 480. A wider pool spreads the same misassigned budget over more candidates, which is consistent with the boundary-concentration mechanism of the earlier stability work.
- **The $\rho = 0$ controls reproduce their ancestors.** 1.4910 against 1.4943 at $(32, 480)$ and 1.5006 against 1.4842 at $(32, 224)$ – -0.003 and $+0.016$ nats on different hardware and a newer torch stack, within this project’s single-seed spread. Together with the bitwise unit test of the $\rho = 0$ code path, the recent codebase changes leave the baseline intact.

The answer to the motivating question is therefore no: the surrogate’s per-neuron credit assignment is its active ingredient. What survives permutation – the budget, the kernel geometry, the zero-sum structure – is worth less than nothing on its own.

2 Setup

Sixteen runs on a fresh 8-GPU machine: $K = 32$, $J \in \{32, 96, 224, 480\}$, $\rho \in \{0, 0.25, 0.5, 1.0\}$, all `through_rank_kappa` at $T = 2$ with no output norm, $b_0 = 0$, full magnitude–direction decoupling, 20k steps. Every run inherits the dumped `config.json` of its ancestor from the earlier campaigns (the stabilization campaign’s $T = 2$ arm at $J = 32$ and 480, the Top- $(K+J)$ campaign’s at $J = 96$ and 224) with only ρ overridden, so architecture, data, optimizer, schedules and seed are identical to the references by construction.

The permutation: per (token, block) row, per backward, $\text{round}(\rho \cdot (K+J))$ positions sampled uniformly without replacement, values shuffled by a uniform permutation of the subset, fresh randomness every step; evaluation is never permuted. The four wide-window cells ($J = 224, 480$) ran to completion; the eight narrow-window cells were stopped by hand at steps 1500–3500 to free the machine once the wide cells had settled the ordering – their curves are reported as observed, and none had been stopped by the divergence guard at that point.

3 Results

Final validation CE after 20k steps (completed cells), with the gap to the unpermuted run of the same cell:

(K, J)	$\rho = 0$	$\rho = 0.25$	$\rho = 0.5$	$\rho = 1$
(32, 224)	1.5006	1.5381 (+0.038)	1.5761 (+0.076)	1.7925 (+0.292)
(32, 480)	1.4910	1.5025 (+0.012)	1.5566 (+0.066)	1.7269 (+0.236)
hard Top-32, no surrogate (reference): 1.6162				

The degradation is monotone in ρ at both windows and visibly ordered from about step 1000 onward. At $\rho = 1$ both cells end 0.11–0.18 nats *above* the no-surrogate baseline. Partial permutation degrades sublinearly at $J = 480$ ($\rho = 0.25$ costs little) and closer to linearly at $J = 224$.

The stopped narrow-window cells, last validation before the stop:

(K, J)	$\rho = 0$	$\rho = 0.25$	$\rho = 0.5$	$\rho = 1$
(32, 32)	1.9615 (2500)	2.0313 (2500)	2.2246 (2000)	7.6565 (1500)
(32, 96)	1.8389 (3500)	1.9070 (3500)	2.0603 (3000)	5.9900 (2500)

The same monotone ordering holds at every observed step, and the $\rho = 1$ cells are not merely behind but diverging – CE 6–8 and rising – where the wide-window $\rho = 1$ cells train stably to the end.

The control. $\rho = 0$ against the ancestor run of the same cell: 1.4910 vs 1.4943 at (32, 480), 1.5006 vs 1.4842 at (32, 224). The first agrees to 0.003 nats; the second differs by +0.016, at the upper end of the run-to-run variation this project has seen from identical configs across machines, and the two curves track each other throughout training. Since the $\rho = 0$ code path is additionally covered by a bitwise identity test, we read both as agreement: the sf, value-gradient, permutation and data-parallel changes to the codebase did not move the baseline.

4 Interpretation

The control was designed so that everything except credit assignment survives: each permuted position still receives a correctly-scaled, correctly-localized kick, each token still injects the same

Permuting a ρ fraction of the surrogate signal dL/dp (rblapsum kappa, $T = 2$): quality is monotone in the intact fraction




Figure 1: Validation CE by window, ρ light to dark, the california ancestor of each cell dashed. In the completed panels (bottom) the $\rho = 0$ curve lies on its ancestor and the ordering in ρ is monotone throughout; in the stopped panels (top) the tick marks the last measurement, with $\rho = 1$ already above the axis range and climbing.

multiset of signal values into the same kernel profile, and the correction still removes the common mode. If the surrogate worked through its budget or its geometry, $\rho = 1$ would approximately reproduce it. Instead $\rho = 1$ is worse than deleting the term entirely, which pins the surrogate’s value – and more – on *which neuron gets which signal*.

That full permutation is worse than nothing admits a mechanical reading consistent with the earlier scale-dynamics measurements, though we mark it as interpretation. The misassigned kicks are refreshed every step, so their first-order effect averages toward zero; what does not average out is second-order: every step executes full-magnitude updates on wrongly-chosen rows (Adam normalizes per parameter, so a persistent stream of large misassigned gradients produces full-size steps), which is pure heating of the score scale plus a per-step scrambling of the ranking the forward depends on. The narrow-window divergences fit the same picture: at $J = 32$ the permuted budget lands on 64 candidates instead of 512, the per-neuron kick is an order of magnitude larger, and the boundary-referenced geometry that the kappa-ablation note showed to be the fragile quantity is kicked correspondingly harder.

For the programme, the positive statement is the useful one: the RBLapSum backward, non-derivative as it is, carries real per-neuron task information – the “wrong gradients” hypothesis is rejected in its strong form. The cheap follow-up, not run, is the same sweep on the boundary term alone (permuting only the $-q \Sigma a$ routing while leaving the direct term intact), which would split the credit between the two components of the score path.

5 Scope

Single runs at one seed, one active count ($K = 32$), one variant ($T = 2$, no output norm), and a coarse ρ grid. The narrow-window cells are truncated at steps 1500–3500 by an operational decision, not by the guard: their divergence-in-progress at $\rho = 1$ is an observation about trajectories, not recorded final outcomes, and earlier work found marginal-cell divergence stochastic. The ancestors were trained on a different machine and torch version, so the control comparison inherits cross-machine variation (± 0.01 – 0.02 nats). The permutation acts before the kernel weighting by design; the harsher post- κ variant (scrambling the scale profile along with the signal) was specified but not run, and the permutation of the boundary routing alone was not attempted.