

The soft-forward bottleneck

Research note: `rblapsum_sf` against its hard-forward twin and the hard Top- K' ladder

2026-09-27.

1 Purpose and summary of conclusions

`rblapsum_sf` keeps RBLapSum’s Top- $(K+J)$ candidate pool, rank boundary $b = \max(b_0, s_{(K+1)})$ and exponential kernel, but places the probabilities in the forward pass: every pool member outputs $z_i p_i$ with $p_i = F((s_i - b)/T)$ (the Laplace CDF), features outside the pool output exactly zero, and the backward pass is the gradient of this forward – there is no train-time discrepancy between the function trained and the function differentiated. The boundary’s own derivative is routed by the same switch as in the hard-forward mode and is κ -weighted (`through_rank_kappa`) throughout this note.

Because the trained forward is soft, two validation metrics exist and they disagree strongly. `val/ce` evaluates the model under the hard Top- K forward that every other regime uses at inference (only the K selected activations participate), and is directly comparable across regimes. `val_soft/ce` evaluates the soft forward actually trained. This note reports both, against two references: the hard-forward `through_rank_kappa` twin of every cell, and the hard Top- K' post-norm baselines at the matching pool $K' = K + J$.

Conclusions:

- **Under its own forward, the soft model beats the hard Top- K' baseline with the same pool in every cell**, by 0.055 to 0.114 nats. The best values are the best in the project so far: 1.3430 at $(K, J) = (64, 448)$ and 1.3498 at $(32, 480)$, both at $T = 2$, against 1.4232 for hard Top-512 and 1.4146 for the best hard-forward surrogate cell.
- **Under the hard Top- K forward, the soft model loses to its hard-forward twin in every cell**, by 0.086 to 1.14 nats, and widening J makes it monotonically worse – the opposite of the hard-forward family, where widening J helps. The damage tracks the probability mass the trained forward carries outside the top K (measured below).
- **The two metrics invert every ordering the earlier campaigns established.** Under soft evaluation: widening J helps monotonically with no reversal at $(32, 480)$; $T = 2$ is better than the post-norm in eight of nine comparable cells (they are equal within 0.001 at $(128, 384)$); and at a fixed pool the active count K barely matters – the spread across K is 0.006–0.020 nats, an order of magnitude smaller than the pool effect.
- **Stability is not improved.** One cell diverged: $(128, 128)$ at $T = 2$, stopped by the guard at step 3480, a cell whose hard-forward twin survives. $(128, 384)$ at $T = 2$ had a large recovered excursion near step 7000. The post-norm arm survived everywhere.
- **As a route to a hard Top- K model, the soft forward currently pays for itself:** the smallest hard-eval penalty in the grid is +0.086 nats at $(256, 256)$ post-norm. As a model evaluated the way it is trained – a weighted pool of $K+J$ features – it is the strongest configuration tested in this project at every pool size.

2 Setup

Twenty runs: the ten (K, J) cells of the Top- $(K+J)$ -versus-hard note ($K, K+J \in \{32, 64, 128, 256, 512\}$, $J = K' - K > 0$), each under the two variants of that note – $T = 2$ with no output norm, and $T = 1$ with an RMSNorm on each bottleneck’s output (“post-norm”). Every run inherits the dumped `config.json` of its hard-forward `through_rank_kappa` twin (the stabilization campaign for the six cells trained there, the Top- $(K+J)$ campaign for the rest), with only the surrogate mode overridden: architecture, data, optimizer, schedules, seed, $b_0 = 0$ and the support scale 1.0 are identical by construction. Divergence guard as in the earlier campaigns.

The two references. The hard-forward twins’ curves are the original runs of their campaigns. The hard Top- K' post-norm baselines are the same five runs used in the Top- $(K+J)$ note. Both were recovered from the Drive backups of the retired machine; the new runs trained on the successor box from the same repo state.

Evaluation. `val/ce` is computed with `hard_inference`: the forward keeps the top K activations and drops the rest, exactly the inference forward of every other regime, so the number is comparable on the same TensorBoard plot. `val_soft/ce` is a second evaluation pass with the trained soft forward, $z_i p_i$ over the pool. The training-side `active_count` (L0) remains the hard support size K ; the pool’s probability mass $\sum_i p_i$ is logged as `rb_soft_mass`.

3 Results

Hard evaluation: the comparable metric

Final `val/ce` after 20k steps, soft-forward (sf) against its hard-forward twin (kappa):

(K, J)	$T = 2$			post-norm		
	kappa	sf	gap	kappa	sf	gap
(32, 32)	1.5492	1.6815	+0.132	1.5310	1.6574	+0.126
(32, 96)	1.5146	1.9369	+0.422	1.5115	1.8368	+0.325
(32, 224)	1.4842	2.2553	+0.771	1.4949	2.0165	+0.522
(32, 480)	1.4943	2.6381	+1.144	1.5370	2.1118	+0.575
(64, 64)	1.5142	1.6450	+0.131	1.4834	1.6005	+0.117
(64, 192)	1.4922	1.8927	+0.401	1.4602	1.7563	+0.296
(64, 448)	1.4727	2.1958	+0.723	1.4538	1.9166	+0.463
(128, 128)	1.4802	diverged	(3480)	1.4423	1.5441	+0.102
(128, 384)	1.4598	1.8537	+0.394	1.4294	1.6954	+0.266
(256, 256)	1.4425	1.8550	+0.413	1.4146	1.5008	+0.086

The soft model loses everywhere under this metric, the loss grows with J at every K , and it shrinks with K at fixed J/K . The post-norm variant is uniformly less damaged than $T = 2$.

Soft evaluation: the trained function

Final `val_soft/ce`, grouped by pool, against hard Top- K' with the same pool (post-norm, hard evaluation – its trained forward):

pool K'	(K, J)	sf $T=2$	sf post-norm	hard Top- K'	best margin
64	(32, 32)	1.4269	1.4511	1.5408	−0.114
128	(32, 96)	1.3852	1.4054	1.4938	−0.109
	(64, 64)	1.3913	1.4039		−0.103
256	(32, 224)	1.3645	1.3931	1.4560	−0.092
	(64, 192)	1.3667	1.3736		−0.089
	(128, 128)	—	1.3762		−0.080
512	(32, 480)	1.3498	1.3846	1.4232	−0.073
	(64, 448)	1.3430	1.3633		−0.080
	(128, 384)	1.3630	1.3624		−0.060
	(256, 256)	1.3571	1.3653		−0.066

Every cell beats the pool-matched hard baseline. Within a pool the spread across K is small (0.006 at pool 128, 0.012 at 256, 0.020 at 512), and the smallest K is never the worst; the pool size dominates. $T = 2$ is ahead of the post-norm in every comparable cell except (128, 384), where the two agree within 0.001.

Where the hard-eval damage comes from

$\text{rb_soft_mass} = \sum_i p_i$ over the pool at the end of training, minus K , is the probability mass the trained forward carries outside the kept set – the part the hard evaluation deletes. Post-norm variant:

(K, J)	mass − K	hard-eval gap to kappa	gap per unit mass
(32, 32)	2.9	+0.126	0.043
(32, 96)	9.3	+0.325	0.035
(32, 224)	18.2	+0.522	0.029
(32, 480)	30.9	+0.575	0.019
(64, 64)	5.0	+0.117	0.023
(64, 192)	18.5	+0.296	0.016
(64, 448)	36.9	+0.463	0.013
(128, 128)	9.6	+0.102	0.011
(128, 384)	33.6	+0.266	0.008
(256, 256)	17.3	+0.086	0.005

Within each K the gap grows monotonically with the tail mass; across K the damage per unit of tail mass falls roughly as $1/K$, consistent with each deleted feature mattering less when more are kept. The $T = 2$ variant carries $1.0\text{--}2.0\times$ more tail mass than the post-norm in every cell (e.g. 92.8 vs 30.9 at (32, 480)), matching both its better soft and worse hard evaluation.

Stability

(128, 128) at $T = 2$ diverged: best CE 1.74 at step ~ 3100 , then a gradient-norm blowup to $\sim 10^{11}$, stopped by the guard at 3480 after 300 steps above CE 7. Its hard-forward twin survives the same cell (final 1.4802), so the soft forward is not a stabilization. (128, 384) at $T = 2$ burst near step 6500–7500 and recovered (visible in the pool-512 figure). All ten post-norm cells and the remaining $T = 2$ cells reached 20k steps.

Figures

Each pool figure: left $T = 2$, right post-norm; the hard-forward family in blues, the soft-forward family in oranges, both light to dark with increasing K ; hard Top- K' dashed black. The vertical

axis is the hard-eval val/ce , so these figures show the comparable metric; the soft-eval figure that follows them shows $\text{val_soft}/\text{ce}$ for the soft family against the same hard ladder.




Figure 1: Pool $K' = 64$. The single cell (32,32): the soft model’s hard-eval curve runs ~ 0.13 nats above its twin in both variants.




Figure 2: Pool $K' = 128$. The $T = 2$ soft cells drift *upward* over the second half of training as the tail mass grows; the post-norm holds flat but offset. The kappa family and the hard baseline are unchanged from the earlier note.

4 Interpretation

The two evaluations measure different models, and the tables say the trained one is good and the hardened one is not. Training optimizes the weighted-pool forward, and the optimizer uses the freedom this buys: it parks measurable probability mass on the J tail features (3–93 per-token units of $\sum p$ beyond K), and the task loss of the trained function improves with every widening of the pool – past the point where the hard-forward family reverses, and below every hard baseline with the same pool. Hard evaluation then deletes exactly that mass, and the measured damage is proportional to it within each K . Nothing here is surprising once stated this way; the quantitative point is that the exchange rate is poor everywhere in the grid: even the best cell pays +0.086 nats for hardening, which is more than the entire margin the hard-forward surrogate holds over hard Top- K' .

The inverted orderings follow from the same picture. A wider kernel ($T = 2$) fattens p in the tail, which is capacity under soft evaluation and liability under hard; the post-norm, which controls

the score scale rather than the kernel, keeps the profile sharper – better hardened, worse soft. At a fixed pool the active count K mostly relabels which features sit above the boundary reference without changing the pool’s capacity, so the soft loss barely moves with K while the hard loss moves strongly.

Read as a design, `rb_lapsum_sf` at these settings is a LapSum-style soft bottleneck with a rank-defined boundary and no barrier solve, trained exactly as evaluated. On that reading it is the best model in the project at every pool size, and the natural follow-ups are the ones that would close the train/inference gap rather than compare it to hard models it does not implement: annealing T toward a hard forward over training, a hardening fine-tune phase, or an explicit penalty on the tail mass (`rb_soft_mass` is already the right observable, and its correlation with the hard-eval gap suggests a target). None of these were run.

5 Scope

All cells are single runs at one seed, and the hard-forward references are reused from the earlier campaigns rather than retrained – identical configs and seed, but a different machine (5090s against 4090s), so bit-level reproduction is not implied. The soft-forward evaluation exists only for the new family; the hard baselines have no soft variant to compare against. Only one boundary mode (`through_rank_kappa`) and support scale (1.0) were trained; the ablations the mode switches were built for – `detach`, `through_rank`, scaled boundary terms – are untouched. The divergence at (128, 128), $T = 2$ is a single cell and earlier work found marginal-cell divergence to be stochastic, so it should be read as a rate, not a verdict. Finally, the soft model’s inference-time cost differs from a hard Top- K model’s: its forward keeps $K+J$ weighted features, so comparisons at matched K favour it unfairly on quality and disfavour it on sparsity; the pool-matched comparison in the second table is the fair one.

×




Figure 3: Pool $K' = 256$. The soft $T = 2$ cell at $(128, 128)$ ends at its guard stop (cross). Hard-eval damage orders inversely with K in both variants.




Figure 4: Pool $K' = 512$. The widest windows: the soft family sits 0.09–1.1 nats above the blue twins under hard eval, with the $(128, 384)$ $T = 2$ excursion near step 7000 visible; the recovered curve returns to its trend.




Figure 5: The soft family under its own forward (*val_soft/ce*), one panel per pool, against the pool-matched hard Top- K' baseline (dashed). Solid $T = 2$, dotted post-norm. Every curve ends below the dashed reference; within a panel the curves cluster tightly across K .