

Stabilizing RBLapSum `through_rank_kappa`

Research note: kernel temperature, learning rate, output normalization and support-gradient scale

2026-09-21/22.

1 Purpose and summary of conclusions

RBLapSum with the kappa correction is usable but not unconditionally stable: at $K = 128$ and $T = 1$ it destabilizes within a few thousand steps. Four interventions are available, and they act on different parts of the mechanism identified earlier (`docs/scale-dynamics-note-neutral.tex`): broadening the kernel ($T = 2$) lowers the surrogate gain, lowering the learning rate shrinks both the signed drift and the second-order heating, an RMSNorm on each bottleneck’s output removes the depth amplification of the score scale, and scaling the support term down attenuates the boundary-exchange gradient directly.

This note compares all four against the untouched baseline on the same six (K, J) configurations, and asks which one stabilizes training and which one converges best.

Conclusions:

- **All four interventions stabilize.** The baseline diverges in the two $K = 128$ configurations, at steps 1680 and 2200. Every other arm reaches 20k steps in all six configurations; no divergence was observed in any of the 24 intervention runs. Stabilization alone therefore does not separate the four – the cost of stabilizing does.
- **The output RMSNorm is the best intervention on both counts.** It converges best in five of six configurations, including both configurations the baseline cannot survive, and its 1.4294 at $(K, J) = (128, 384)$ is the best value in the comparison. It is also the only arm that stabilizes without a measurable accuracy cost.
- **Broadening the kernel is a close second** and needs no architectural change: within ± 0.043 nats of the baseline wherever the baseline survives, and 0.030–0.038 nats behind the RMSNorm in the $K = 128$ rows it rescues.
- **Lowering the learning rate stabilizes by training less.** It is last in all six configurations, 0.080–0.155 nats behind the RMSNorm.
- **Scaling the support gradient to 0.3 stabilizes reliably but costs accuracy.** It never diverges, but it is 0.005–0.063 nats behind the RMSNorm and is worse than the untouched baseline in every configuration where the baseline survives.
- **The RMSNorm also helps the hard bottleneck**, which has no surrogate gradient and no instability to fix: 1.6022 against 1.6162 at $K = 32$ and 1.4232 against 1.4298 at $K = 512$. It is therefore not a surrogate-specific crutch.

2 Setup

Thirty runs, all RBLapSum `through_rank_kappa` under full magnitude–direction decoupling, with the architecture, data, optimizer, schedules and seed of the earlier campaigns (8×768 , `abs_topk` over $N = 1536$ features, `residual_out` placement, batch 96×512 tokens, 20k steps, $b_0 = 0$).

Six configurations,

$$(K, J) \in \{(32, 32), (32, 480), (64, 64), (64, 448), (128, 128), (128, 384)\},$$

each under five arms:

arm	T	lr	change
baseline	1	6×10^{-4}	—
$T = 2$	2	6×10^{-4}	wider kernel
lr 2×10^{-4}	1	2×10^{-4}	smaller steps
post-norm	1	6×10^{-4}	RMSNorm on each bottleneck’s output
support grad 0.3	1	6×10^{-4}	<code>rblapsum_support_scale=0.3</code>

The first four arms ran on one 8-GPU machine (24 runs, names `ca_rout_rblapsum_kappa*_stab_*`); the support-gradient-scale arm ran on a second machine (6 runs, `ko_rout_rblapsum_kappa*_supp03_*`), along with two hard-AbsTopK runs used for the last conclusion above.

Two implementation notes. The output RMSNorm (`activation_bottleneck.post_norm`) is applied inside the bottleneck module, and the bottleneck whose output already feeds a norm – `residual_out` in the final block, ahead of `norm_f` – is exempt. The support-gradient scale multiplies the boundary-exchange term g_s *after* the mode correction, so the zero-sum structure of `through_rank_kappa` is preserved and the term is uniformly rescaled; the hard task path is untouched, so the scale interpolates between the full surrogate (1.0) and the plain hard-TopK backward (0.0).

A run is stopped and recorded as diverged when its training CE becomes non-finite, stays above its own best plus 2.5 nats for 800 consecutive steps, or stays above 7.0 for 300 steps after step 1500. The window is relative to each run’s own best, so the slower low-learning-rate arm is not penalized for being behind.

3 Results

(K, J)	baseline	$T = 2$	lr 2×10^{-4}	post-norm	support grad 0.3
(32, 32)	1.5521	1.5492	1.6709	1.5310	1.5847
(32, 480)	1.4900	1.4943	1.6168	1.5370	1.5419
(64, 64)	1.5281	1.5142	1.6376	1.4834	1.5348
(64, 448)	1.4661	1.4727	1.6006	1.4538	1.5092
(128, 128)	diverged (1680)	1.4802	1.5970	1.4423	1.5037
(128, 384)	diverged (2200)	1.4598	1.5843	1.4294	1.4926

Final validation CE after 20k steps, or the step at which the divergence guard stopped the run. Best value per configuration in bold.

A recovered excursion. The baseline at (64, 64) survives but is visibly marginal. Its validation CE rises from 1.74 to 5.12 at step 6500 and recovers within 500 steps, rises again to 5.60 at step 9000, and decays back below 1.8 by step 11000, ending at its own best value, 1.5281. No other surviving run behaves this way: the only other runs above CE 2.0 after step 3000 are the two




Figure 1: $K = 32$: smaller candidate window on the left, larger on the right; the lower row repeats each panel with the vertical axis capped at 2.0 nats. Colour encodes the arm and line style its kind: the untouched baseline is dashed, the two arms that stabilize without an accuracy cost ($T = 2$ and the output RMSNorm) are solid, and the two attenuating arms (reduced learning rate, support-gradient scale) are dash-dotted and slightly transparent. Every arm survives here, so this configuration measures cost rather than stabilization: the baseline and $T = 2$ curves are indistinguishable, the output RMSNorm wins at $J = 32$ but loses 0.047 nats at $J = 480$, and the reduced learning rate is clearly behind throughout.

slowest reduced-learning-rate runs, which are simply behind rather than excusing. This is the burst-and-recover behaviour the divergence guard’s 800-step window is meant to tolerate, and it places the baseline at the edge of stability one configuration earlier than the divergences alone would suggest.

Where the arms differ. Relative to the output RMSNorm, aggregated over the configurations in which both runs converged:

arm	gap to post-norm (nats)	configurations
baseline	−0.047 to +0.045	4 (diverges in 2)
$T = 2$	−0.043 to +0.038	6
support grad 0.3	+0.005 to +0.063	6
$\text{lr } 2 \times 10^{-4}$	+0.080 to +0.155	6

The baseline and $T = 2$ straddle zero: they beat the RMSNorm at (32, 480) by 0.047 and 0.043 nats respectively and lose to it everywhere else. The other two arms are behind in every configuration.

The hard bottleneck. Two hard-AbsTopK runs with the same output RMSNorm test whether the norm helps where there is no surrogate gradient and no instability. Against the matched runs from the earlier campaign (same configuration and seed, no norm):




Figure 2: $K = 64$, panels as above. No run diverges, but the baseline (blue, $J = 64$) leaves the plotted range twice and returns – the recovered excursion discussed in the text. The output RMSNorm is best at both J , by 0.012–0.045 nats over the baseline; the capped row makes the ordering of the four arms legible over the second half of training.

K	without post-norm	with post-norm	difference
32	1.6162	1.6022	−0.0140
512	1.4298	1.4232	−0.0066

Both improve. The norm is therefore not compensating for the surrogate; it improves a bottleneck that has no surrogate at all.

4 Interpretation

The ordering is consistent with the initialization measurement in `docs/init-pi-grid.pdf` and `docs/init-pi-grid-postnorm.pdf`. Under `residual_out` each block replaces the residual stream, so the bottleneck’s output scale compounds through depth: at $K = 512$ the pre-gate score RMS runs $1.67 \rightarrow 6.16 \rightarrow 13.3$ across blocks 0, 4 and 7 at initialization, and the surrogate gain Π with it, $262 \rightarrow 1089 \rightarrow 2384$. Adding the output RMSNorm holds the score RMS at $1.67 \rightarrow 1.79 \rightarrow 1.82$ and Π at $262 \rightarrow 288 \rightarrow 293$: the deep-block gain falls by a factor of eight and becomes depth-uniform.

That is a different kind of fix from the other three. Broadening the kernel, shrinking the step and damping the support term all reduce the *size of the surrogate contribution*; the norm removes the *amplification* that made the contribution dangerous in deep blocks, leaving the gradient itself intact. The measured costs follow: the three attenuating arms pay between 0.03 and 0.16 nats, while the norm pays nothing and usually gains.

The support-gradient-scale result is worth stating explicitly, because it is the intervention the earlier counterfactual experiments identified as causally effective: disabling the support gradient




Figure 3: $K = 128$, panels as above. The baseline (blue) diverges at both J and leaves the plotted range, its stop marked by a cross at the top edge of each panel. All four interventions converge, and the ordering is the same at both J and stable from about step 3000: the output RMSNorm (red) runs lowest throughout, ahead of $T = 2$, then the support-gradient scale, then the reduced learning rate.

stopped both burst initiation and post-collapse inflation. It does stabilize here – six configurations, no divergences – but at 0.3 it costs 0.005–0.063 nats relative to the norm and is worse than the untouched baseline wherever the baseline survives. Attenuating the boundary exchange removes the instability and part of the learning signal together.

5 Recommendation and scope

For the configurations tested, the output RMSNorm is the intervention to prefer: it is the only one that stabilizes the $K = 128$ rows without paying for it, it is best in five of six configurations, and it also improves the hard bottleneck. Where an architectural change is unwanted, $T = 2$ is the next best option and costs little.

Scope. All cells are single runs at one seed, and earlier work in this project showed that divergence at marginal configurations is stochastic, so individual outcomes should be read as draws from a divergence rate rather than as deterministic verdicts; the patterns that repeat across configurations – the baseline being the only arm that diverges, the RMSNorm leading in five of six, the reduced learning rate being uniformly last – are the robust part. Only one support-gradient-scale value (0.3) and one reduced learning rate (2×10^{-4}) were tested, so neither arm is tuned; a milder support-gradient scale in particular might retain the stabilization at a smaller accuracy cost. Finally, the single configuration in which the untouched baseline is best, (32, 480), is also the least stressed one, so the norm should be read as a fix for stressed configurations rather than a universal improvement.